

# TEAM 11 – VINYL RECORD AND BOOK CLASSIFIER

PROJECT 3: MARK WIREMAN  
APRIL 10, 2025

# PROBLEM STATEMENT

With a growing collection of books and vinyl albums it is a challenge to keep track of what I have, the vintage of the item, and current value.

Alas, a solution to the problem has been found! What a better way to use my newly acquired knowledge to begin a journey along the AI deep learning path to assist me in creating an inventory of my collection.

The first step is to get the data, then train the model....



# IMAGE CLASSIFICATION OVERVIEW

## AGENTIC AI SYSTEM



### DOWNLOAD IMAGES

Uses the DuckDuckGo Python library to download random images of books and vinyl records.



### TRAIN THE MODEL

Designed to train and validate a deep learning model using either MobileNetV2 or ResNet50 as the base model



### PREPROCESS IMAGES

Perform a cleaning of the images using the Tenserflow and PIC libraries with actions to convert or remove, if required.



### PREDICT THE CATEGORY

The model is intended to classify images into two categories: vinyl records and books.



### CHOOSE THE PARAMETERS

A Flask application to allow the user to configure the agents and select what agents to run.




### VIEW RESULTS

A built-in system to view progress in real-time and visualize the results.

# KEY FEATURES

- **Model Selection:** The code allows users to choose between MobileNetV2 and ResNet50 as the base model.
- **Custom Preprocessing Layer:** A custom preprocessing layer is implemented to handle image preprocessing for both MobileNetV2 and ResNet50.GPU
- **Optimization:** The code is optimized for GPU usage with shared memory optimization and auto-clustering enabled.
- **XLA Compilation:** The code uses XLA (Accelerated Linear Algebra) compilation to optimize computations.
- **K-Fold Cross-Validation:** The code implements K-Fold cross-validation to evaluate the model's performance.
- **Data Augmentation:** The code applies data augmentation techniques to the training data.
- **Model Evaluation:** The code evaluates the model's performance using metrics such as accuracy, precision, and recall.








# PERFORMANCE AND BENCHMARKING FEATURES

LET THE AGENTS DO THE WORK!

# AUTOMATED PERFORMANCE AND BENCHMARKING

- **GPU Performance Benchmarking:** The code includes a benchmarking function to measure the performance of the GPU with different batch sizes.
- **Batch Size Optimization:** The code adjusts the batch size based on the GPU's performance to optimize shared memory utilization.
- **Speedup Calculation:** The code calculates the speedup achieved by using batch processing with different batch sizes.
- **Custom Preprocessing Layer:** A custom preprocessing layer to handle image preprocessing
- **Global Average Pooling Layer:** A global average pooling layer to reduce the spatial dimensions of the feature maps
- **Dense Layers:** One or more dense layers with ReLU activation and dropout
- **Final Classification Layer:** A final classification layer with softmax activation



# AUTOMATED PERFORMANCE AND BENCHMARKING

```
Cross-Validation Results:  
Average Validation Loss: 0.2789  
Average Validation Accuracy: 0.9164  
Best model from fold 12 with validation loss: 0.1443  
  
Training final model on all training data...  
Found 226 images belonging to 2 classes.
```

Metrics from folds used  
to determine “best  
model” to use

```
# Create a new model for this fold  
fold_model = create_model(base_model_name=model,no_of_layers=layers,use_xla=use_xla)  
  
# Setup callbacks  
fold_callbacks = []  
  
# Add early stopping  
early_stopping = tf.keras.callbacks.EarlyStopping(  
    monitor='val_loss',  
    patience=10,  
    restore_best_weights=True  
)  
fold_callbacks.append(early_stopping)  
  
# Add learning rate scheduler  
reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(  
    monitor='val_loss',  
    factor=0.5,  
    patience=5,  
    min_lr=1e-6,  
    verbose=1  
)  
fold_callbacks.append(reduce_lr)
```

Configuring  
each fold for  
early stopping  
and learning  
rate

```
# benchmark function to measure performance improvements with shared memory  
def benchmark_gpu_performance(model, img_paths, runs=5, batch_sizes=[1, 4, 8, 16, 32, 64]):  
    """Benchmark GPU performance with different batch sizes to show shared memory benefits"""  
    if not img_paths:  
        print("No images provided for benchmarking")  
        return  
  
    # Limit to first 100 images for benchmarking  
    test_paths = img_paths[:100] if len(img_paths) > 100 else img_paths  
  
    print("\n== GPU SHARED MEMORY PERFORMANCE BENCHMARK ==")  
    print(f"Testing with {len(test_paths)} images, ...")  
  
    results = {}  
  
    # Test individual image processing (no shared memory benefit)  
    print("\nBenchmarking individual image processing (baseline)...")  
    start_time = time.time()  
    for _ in range(runs):  
        for path in test_paths[:10]: # Limit to 10 images for individual processing  
            _ = predict_image(model, path)  
    individual_time = (time.time() - start_time) / (10 * runs)  
    print(f"Average time per image (individual): {individual_time*1000:.2f} ms")  
    results['individual'] = individual_time * 1000  
  
    # Test batch processing with different batch sizes  
    for batch_size in batch_sizes:  
        print(f"\nBenchmarking batch size {batch_size}...")  
        start_time = time.time()  
        for _ in range(runs):  
            _ = predict_batch(model, test_paths, batch_size=batch_size)  
        batch_time = (time.time() - start_time) / (len(test_paths) * runs)  
        print(f"Average time per image (batch size {batch_size}): {batch_time*1000:.2f} ms")  
        results[f'batch_{batch_size}'] = batch_time * 1000  
  
    # Calculate speedup  
    speedup = individual_time / batch_time  
    print(f"Speedup with batch size {batch_size}: {speedup:.2f}x")  
    results[f'speedup_{batch_size}'] = speedup  
  
    # Find the best batch size  
    best_batch_size = batch_sizes[0]
```

Use GPU to calculate best  
batch size to use

```
Epoch 33/100  
26/26 ----- 0s 520ms/step - accuracy: 0.8950 - loss: 0.3043 - precision: 0.8950 - recall: 0.8950  
Epoch 33: val_loss did not improve from 0.33897  
Epoch 33: loss=0.3371, acc=0.8792  
26/26 ----- 15s 579ms/step - accuracy: 0.8944 - loss: 0.3055 - precision: 0.8944 - recall: 0.8944 - val_accuracy: 0.8947 - val_loss: 0.3884 - val_precision: 0.8947 - val_rec  
all: 0.8947 - learning_rate: 2.5000e-05  
Epoch 34/100  
26/26 ----- 0s 520ms/step - accuracy: 0.8024 - loss: 0.3996 - precision: 0.8024 - recall: 0.8024  
Epoch 34: val_loss did not improve from 0.33897  
Epoch 34: loss=0.3537, acc=0.8309  
26/26 ----- 16s 609ms/step - accuracy: 0.8035 - loss: 0.3979 - precision: 0.8035 - recall: 0.8035 - val_accuracy: 0.8421 - val_loss: 0.4027 - val_precision: 0.8421 - val_rec  
all: 0.8421 - learning_rate: 2.5000e-05  
Epoch 35/100  
26/26 ----- 0s 521ms/step - accuracy: 0.8872 - loss: 0.3658 - precision: 0.8872 - recall: 0.8872  
Epoch 35: ReduceLROnPlateau reducing learning rate to 1.249999968422344e-05.  
Epoch 35: val_loss did not improve from 0.33897  
Epoch 35: loss=0.3507, acc=0.8792  
26/26 ----- 16s 581ms/step - accuracy: 0.8869 - loss: 0.3652 - precision: 0.8869 - recall: 0.8869 - val_accuracy: 0.8947 - val_loss: 0.3929 - val_precision: 0.8947 - val_rec  
all: 0.8947 - learning_rate: 2.5000e-05  
Fold 7 - Validation Loss: 0.3390, Validation Accuracy: 0.8947
```

ReduceLROnPlateau callback monitors validation  
loss and the training uses other metrics to  
improve model performance and validation.

```
Training fold 8/12  
Found 207 images belonging to 2 classes.  
Found 19 images belonging to 2 classes.  
Created Layer 0 with units of 1024.  
Created Layer 1 with units of 512.  
Created Layer 2 with units of 256.  
Created Layer 3 with units of 256.  
Created Layer 4 with units of 1024.  
Created Layer 5 with units of 512.  
Created Layer 6 with units of 256.  
Created Layer 7 with units of 256.  
Created Layer 8 with units of 1024.  
Created Layer 9 with units of 512.  
Created Layer 10 with units of 256.  
Created Layer 11 with units of 256.  
Created Layer 12 with units of 1024.  
Created Layer 13 with units of 512.  
Created Layer 14 with units of 256.  
resnet_xla.True  
Epoch 1/100
```

K-Fold cross-validation to evaluate the model's  
performance and automated assignment of  
units to the layers to improve the deep learning  
with the training set

```
26/26 ----- 0s 564ms/step - accuracy: 0.4910 - loss: 0.6957 - precision: 0.4910 - recall: 0.4910  
Epoch 1: val_loss improved from inf to 0.69340, saving model to tmp/kfold_20250408-071034/fold_8/checkpoint/best_model.keras  
Epoch 1: loss=0.6942, acc=0.5072  
26/26 ----- 31s 841ms/step - accuracy: 0.4916 - loss: 0.6956 - precision: 0.4916 - recall: 0.4916 - val_accuracy: 0.4737 - val_loss: 0.6934 - val_precision: 0.4737 - val_rec  
all: 0.4737 - learning_rate: 1.0000e-04  
Epoch 2/100
```



# THE FUTURE

BACKLOG OF ENHANCEMENTS



# FUTURE ENHANCEMENTS

1. Allow user to enter a learning rate
2. Allow user to enter a patience value
3. Allow user to enter the number of training images to fetch
4. Create a CSV or database inventory of the predicted images

